

# Artificial Intelligence and Machine Learning

Christian Cabrera Jojoa

Senior Research Associate and Affiliated Lecturer

Department of Computer Science and Technology


University of Cambridge

[chc79@cam.ac.uk](mailto:chc79@cam.ac.uk)

# The ML Context

# Artificial Intelligence (AI)

"The field of **Artificial Intelligence (AI)** is concerned with *understanding and building intelligent entities*."



Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach (4th Edition)*

# AI Dimensions

Human

Rational

Thought

Behaviour

# AI Dimensions

|           | Human                                                                                                                                                                                                                  | Rational                                                                                                                                                                                                                      |
|-----------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Thought   | <p>Systems that <i>think like humans</i>, focusing on cognitive and mental models that <i>mimic human reasoning</i>.</p> <ul style="list-style-type: none"><li>• Example: Context-aware chatbots</li></ul>             | <p>Systems that <i>think rationally</i>, using logical principles and formal reasoning to <i>solve problems and make decisions</i>.</p> <ul style="list-style-type: none"><li>• Example: Rule-based problem solvers</li></ul> |
| Behaviour | <p>Systems that <i>act like humans</i>, exhibiting behaviours and interactions that are <i>natural and intuitive for human users</i>.</p> <ul style="list-style-type: none"><li>• Example: Human-like robots</li></ul> | <p>Systems that <i>act rationally</i>, making optimal decisions based on available <i>information and defined goals</i>.</p> <ul style="list-style-type: none"><li>• Example: Optimal decision makers</li></ul>               |

# AI Dimensions

|           | Human                                                                                                                                                                         | Rational                                                                                                                                                    |
|-----------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Thought   | <ul style="list-style-type: none"><li>• Cognitive modeling</li><li>• Human-like reasoning</li><li>• Natural language understanding</li><li>• Common sense reasoning</li></ul> | <ul style="list-style-type: none"><li>• Logic-based systems</li><li>• Theorem proving</li><li>• Knowledge representation</li><li>• Expert systems</li></ul> |
| Behaviour | <ul style="list-style-type: none"><li>• Natural language processing</li><li>• Computer vision</li><li>• Robotics</li><li>• Human-computer interaction</li></ul>               | <ul style="list-style-type: none"><li>• Optimization algorithms</li><li>• Game theory</li><li>• Planning systems</li><li>• Decision theory</li></ul>        |

# AI Foundations

## AI Foundations

Philosophy Mathematics Economics Neuroscience

Psychology Computer Engineering Control Theory

Linguistics Mind-body problem Logic Reasoning Ethics Probability


Statistics Optimization Graph theory Decision theory Game theory

Neural networks Learning Cognition Memory Algorithms Systems

Feedback Language NLP Communication




# AI History




# Artificial Intelligence - Intelligent Entities Rows

The field of **Artificial Intelligence** (AI) is concerned with *understanding and building intelligent entities*.



*Stuart Russell and Peter Norvig, Artificial Intelligence: A Modern Approach (4th Edition)*


# Artificial Intelligence - Intelligent Entities Rows and Columns



Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach* (4th Edition)

The field of **Artificial Intelligence** (AI) is concerned with *understanding and building intelligent entities*.

The field of **Artificial Intelligence** (AI) is concerned with *understanding and building intelligent entities*.




Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach* (4th Edition)

# Code Example

```
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
# Generate sample data
X, y = make_blobs(n_samples=100, centers=2, random_state=42)
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()
```

# Code Example in Columns




Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach* (4th Edition)

```
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
# Generate sample data
X, y = make_blobs(n_samples=100, centers=2, random_state=42)
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()

from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
# Generate sample data
X, y = make_blobs(n_samples=100, centers=2, random_state=42)
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()

from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
# Generate sample data
X, y = make_blobs(n_samples=100, centers=2, random_state=42)
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()

from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
# Generate sample data
X, y = make_blobs(n_samples=100, centers=2, random_state=42)
# Plot the data
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Sample ML Dataset")
plt.show()
```



Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach* (4th Edition)

# History of Artificial Intelligence

The field of Artificial Intelligence has evolved through several key periods:

- 1943: McCulloch and Pitts proposed a model of artificial neurons - 1950: Alan Turing published "Computing Machinery and Intelligence" introducing the Turing Test - 1956: The Dartmouth Conference, where the term "Artificial Intelligence" was coined

## The Early Years (1940s-1950s)

- Initial enthusiasm gave way to disappointment - Limited computing power and data availability
- Overpromising and underdelivering on AI capabilities

## **The First AI Winter (1974-1980)**



- Rise of rule-based systems - Success in specific domains like medical diagnosis - Commercial applications of AI began to emerge

## **The Expert Systems Era (1980s)**

- Expert systems proved expensive to maintain - Limited scalability of early AI approaches -
- Reduced funding and interest in AI research

## **The Second AI Winter (Late 1980s-1990s)**

- Machine Learning revolution - Deep Learning breakthroughs - Big Data and increased computing power - Widespread practical applications - Integration into everyday technology

## **The Modern Era (2000s-Present)**

- 1997: IBM's Deep Blue defeats chess champion Garry Kasparov - 2011: IBM's Watson wins Jeopardy! - 2012: AlexNet revolutionizes computer vision - 2016: AlphaGo defeats world champion Lee Sedol - 2020s: Large Language Models and Generative AI

This historical context helps us understand how AI has evolved from theoretical concepts to practical applications that impact our daily lives.

## **Key Milestones**

# Public Perception of AI

The public's understanding and perception of Artificial Intelligence has evolved significantly over time:

- AI as human-like robots (Hollywood effect) - AI having consciousness or emotions - AI being either completely benevolent or malevolent - AI being a single, unified technology

## Common Misconceptions

- Science fiction shaping public expectations - Sensationalist headlines about AI capabilities - Fear-mongering about job displacement - Overemphasis on existential risks

## Media Influence

- AI systems are specialized tools, not general intelligence
- Current AI lacks consciousness or self-awareness
- AI systems reflect their training data and human biases
- AI is a collection of different technologies and approaches

## Reality vs. Fiction



- Job displacement and economic impact - Privacy and surveillance - Bias and discrimination -  
Decision-making transparency - Security and misuse

## Public Concerns

- Potential for medical breakthroughs - Environmental monitoring and protection - Educational assistance - Accessibility improvements - Scientific research acceleration

Understanding these perceptions is crucial for:

- Setting realistic expectations
- Addressing legitimate concerns
- Promoting responsible AI development
- Fostering public trust in AI technologies

## Positive Perceptions

# The ML Concept

# What is Machine Learning?

Machine Learning (ML) is a subset of Artificial Intelligence that focuses on developing systems that can learn from and make decisions based on data.

- A system's ability to automatically learn and improve from experience - Without being explicitly programmed - Through pattern recognition and statistical inference

## Core Definition

1. **\*\*Data\*\*** - Training data - Validation data - Test data - Features and labels

2. **Model**

- Mathematical representation
- Learning algorithm
- Parameters and weights

## Key Components

3. **Learning Process**

- Pattern recognition
- Parameter optimization
- Performance evaluation

1. **\*\*Supervised Learning\*\*** - Learning from labeled data - Classification and regression -  
Examples: spam detection, price prediction

2. **Unsupervised Learning**

- Finding patterns in unlabeled data
- Clustering and dimensionality reduction
- Examples: customer segmentation, anomaly detection

3. **Reinforcement Learning**

- Learning through trial and error
- Reward-based learning
- Examples: game playing, robotics

## Types of Machine Learning

- Feature engineering - Model training - Overfitting and underfitting - Bias-variance tradeoff - Generalization - Evaluation metrics

## Fundamental Concepts



- Handling complex patterns - Adapting to new data - Automating decision-making - Extracting insights from large datasets - Solving problems where traditional programming is impractical

## Why Machine Learning?

# History of Machine Learning

Machine Learning has its roots in several disciplines and has evolved through distinct phases:

- 1943: McCulloch and Pitts' artificial neuron model - 1950: Alan Turing's learning machine concept - 1952: Arthur Samuel's checkers-playing program

## Early Foundations (1940s-1950s)

- 1957: Frank Rosenblatt's Perceptron - Early neural network models - Initial optimism followed by limitations

## The Perceptron Era (1950s-1960s)

- Reduced interest in ML research - Focus on expert systems - Limited computational resources

## **The AI Winter Impact (1970s-1980s)**

- Backpropagation algorithm - Support Vector Machines - Decision Trees and Random Forests - Increased computational power

## **The Renaissance (1980s-1990s)**

- Big Data revolution - Deep Learning resurgence - Cloud computing enabling large-scale ML - Open-source ML frameworks

## **The Modern Era (2000s-Present)**

- 1995: Support Vector Machines - 2006: Deep Learning renaissance - 2012: AlexNet and ImageNet success - 2014: Generative Adversarial Networks - 2017: Transformer architecture - 2020s: Large Language Models

## Key Breakthroughs



- Transfer Learning - Few-shot Learning - Self-supervised Learning - Federated Learning - Explainable AI - Ethical AI considerations

This historical perspective shows how ML has grown from theoretical concepts to practical tools that power modern technology.

## Current Trends

# Perception of Machine Learning

Machine Learning is perceived differently across various sectors and communities:

- Competitive advantage - Cost reduction - Process automation - Data-driven decision making - Innovation enabler

## Industry Perspective

- Research frontier - Interdisciplinary field - Theoretical foundations - Experimental validation -  
Peer-reviewed progress

## **Academic View**

- Often conflated with AI - Seen as "magic" or "black box" - Concerns about job displacement - Privacy considerations - Trust and transparency issues

## Public Understanding

- Strategic investment - Digital transformation - ROI considerations - Talent acquisition - Risk management

**Business Leaders**

- Tool for problem-solving - Engineering challenges - Best practices - Open-source collaboration - Continuous learning

## Technical Community

- Bias and fairness - Accountability - Transparency - Privacy protection - Social impact

## Ethical Considerations



- ML can solve any problem
- More data always means better results
- ML models are objective
- ML replaces human judgment
- ML is only for large companies

## Common Misconceptions

- ML as a tool, not a solution - Requires quality data - Needs domain expertise - Involves trial and error - Demands maintenance

Understanding these different perceptions helps in:

- Setting appropriate expectations
- Communicating effectively
- **Realistic Expectations**
- Building trust
- Managing implementation
- Addressing concerns

# The Objective Definition of ML

# Subjective Aspects of Machine Learning

Machine Learning involves several subjective elements that influence its development and application:

- Domain expertise - Intuition in feature selection - Judgment in model evaluation - Interpretation of results - Decision-making thresholds

## Human Factors

- Problem framing - Data collection methods - Feature engineering - Model selection -  
Hyperparameter tuning

## Design Choices

- Value judgments in data selection - Fairness definitions - Privacy boundaries - Transparency requirements - Accountability standards

## Ethical Considerations

- Regional data availability - Language and cultural context - Social norms and values - Regulatory environments - Market expectations

## Cultural Influences



- Confirmation bias - Selection bias - Interpretation bias - Anchoring effects - Overconfidence

## Personal Biases

- Success metrics selection - Performance thresholds - Error tolerance levels - Trade-off decisions
- Quality assessment

## Subjectivity in Evaluation

- Research priorities - Funding decisions - Team composition - Project timelines - Resource allocation

## Impact on Development

- Documentation - Peer review - Diverse perspectives - Continuous feedback - Regular audits

Understanding these subjective elements is crucial for:

- Building better ML systems
- Ensuring fairness
- **Managing Subjectivity**
- Making informed decisions
- Building trust

# Objective Aspects of Machine Learning

Machine Learning is grounded in several objective, measurable components:

- Linear algebra - Probability theory - Statistics - Calculus - Optimization

## Mathematical Foundations

- Algorithms - Data structures - Computational complexity - Memory requirements - Processing speed

## Technical Components

- Accuracy - Precision - Recall - F1 score - ROC curves - AUC scores

## Performance Metrics



- Dataset size - Feature dimensions - Data types - Missing values - Outliers

## Data Characteristics

- Processing power - Memory capacity - Storage requirements - Network bandwidth - Energy consumption

## Computational Resources

- Number of parameters - Model architecture - Training time - Inference speed - Memory footprint

## Model Properties

- Error rates - Loss functions - Convergence speed - Generalization ability - Robustness

## Quality Measures

- Data preprocessing - Model training - Validation procedures - Testing protocols - Deployment pipelines

## Standardized Processes

- Reproducibility - Benchmarking - A/B testing - Statistical significance - Performance baselines

Understanding these objective aspects is essential for:

- Building reliable systems
- Making fair comparisons
- **Objective Evaluation**  
Ensuring reproducibility
- Optimizing performance
- Scaling solutions

# ML Promises

# Promises of Machine Learning

Machine Learning offers numerous potential benefits across various domains:



- Early disease detection - Personalized medicine - Drug discovery - Medical image analysis - Treatment optimization

## Healthcare

- Process automation - Predictive maintenance - Supply chain optimization - Customer behavior analysis - Fraud detection

## **Business and Industry**

- Data analysis acceleration - Pattern discovery - Hypothesis generation - Simulation and modeling - Cross-disciplinary insights

## Scientific Research

- Climate modeling - Resource management - Pollution monitoring - Wildlife conservation -  
Disaster prediction

## **Environmental Protection**

- Personalized learning - Adaptive testing - Educational content creation - Student performance analysis - Learning path optimization

## Education

- Accessibility improvements - Language translation - Disaster response - Public health monitoring - Social service optimization

## **Social Impact**

- Productivity gains - Cost reduction - New business models - Market efficiency - Job creation in new sectors

## **Economic Benefits**

- Autonomous systems - Natural language processing - Computer vision - Robotics - Smart infrastructure

## Technological Advancement



- Personal assistants - Smart homes - Transportation optimization - Healthcare accessibility - Information accessibility

## Quality of Life

- Accelerated innovation - Cross-domain applications - New methodologies - Collaborative research - Knowledge discovery

These promises must be balanced with:

- Ethical considerations
- Privacy concerns
- **Research and Development**
- Implementation challenges
- Resource requirements
- Social impact

# ML Risks

# Risks and Challenges in Machine Learning

Machine Learning implementation carries several significant risks that must be carefully managed:

- Model failure - Data quality issues - System vulnerabilities - Performance degradation - Integration challenges

## Technical Risks

- Bias and discrimination - Privacy violations - Lack of transparency - Accountability gaps - Unintended consequences

## **Ethical Risks**

- Job displacement - Digital divide - Social manipulation - Cultural impact - Dependency creation

## Social Risks

- Adversarial attacks - Data breaches - Model theft - System manipulation - Privacy violations

## Security Risks



- Implementation costs - ROI uncertainty - Market disruption - Resource allocation - Competitive pressure

## Economic Risks

- Regulatory compliance - Liability issues - Intellectual property - Data protection - Contractual obligations

## Legal Risks

- System reliability - Maintenance requirements - Scalability issues - Technical debt - Knowledge gaps

## Operational Risks

- Energy consumption - Resource utilization - E-waste generation - Carbon footprint - Sustainability impact

## Environmental Risks

- Public perception - Trust erosion - Brand damage - Stakeholder confidence - Market position

## Reputational Risks

- Robust testing - Regular audits - Ethical guidelines - Security protocols - Continuous monitoring

Addressing these risks requires:

- Comprehensive planning
- Stakeholder engagement
- **Mitigation Strategies**
- Contingency planning
- Ongoing evaluation

# ML in Practice

# Machine Learning in Practice

Implementing Machine Learning requires careful consideration of practical aspects:



1. **\*\*Problem Definition\*\*** - Business understanding - Success criteria - Resource assessment - Timeline planning - Stakeholder alignment

## 2. Data Preparation

- Data collection
- Cleaning and preprocessing
- Feature engineering
- Data validation
- Version control

## 3. Model Development

- Algorithm selection
- Model training

.. . . .

- Version control - Documentation - Testing - Monitoring - Continuous integration

## Best Practices

- Python ecosystem - TensorFlow/PyTorch - Scikit-learn - Cloud platforms - MLOps tools

## Tools and Frameworks

- Cross-functional teams - Clear communication - Knowledge sharing - Code review - Documentation standards

## **Team Collaboration**

- Model validation - Performance testing - Security testing - Compliance checking - User acceptance testing

## Quality Assurance

- Model retraining - Performance monitoring - Bug fixing - Feature updates - Documentation updates

## Maintenance

- Infrastructure needs - Cost management - Performance optimization - Resource allocation - Load balancing

## Scaling Considerations

- Clear objectives - Quality data - Skilled team - Proper tools - Stakeholder support

## Success Factors



- Data quality - Model drift - Resource constraints - Integration issues - Maintenance overhead

## Common Challenges

- Technical expertise - Project management - Domain knowledge - Resource allocation - Continuous improvement

**Implementing ML successfully requires:**

# Many Thanks!

[chc79@cam.ac.uk](mailto:chc79@cam.ac.uk)